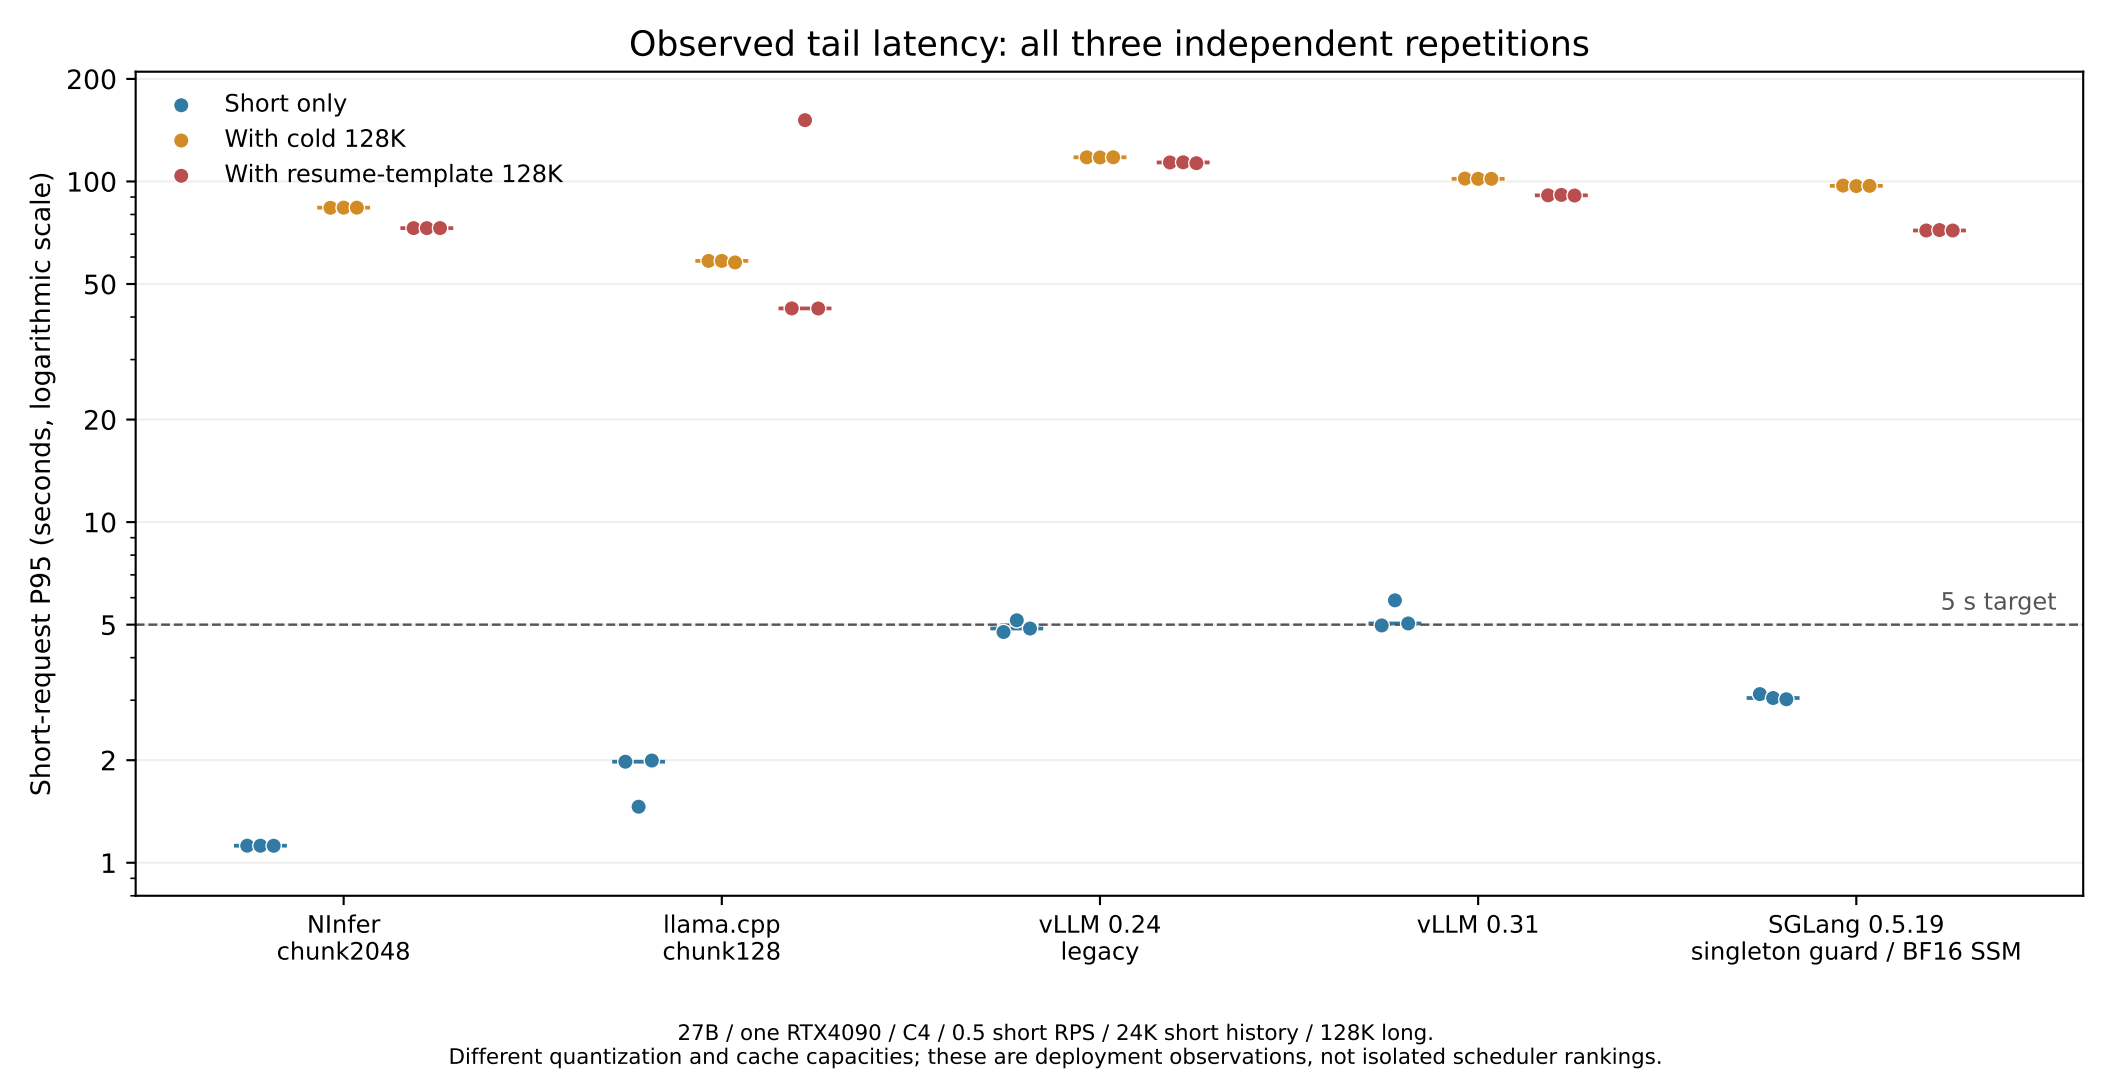
Observed tail latency: all three independent repetitions
Short-request P95 (seconds, logarithmic scale)
200
100
50
20
10
5
2
1
Short only
With cold 128K
With resume-template 128K
5 s target
NInfer
chunk2048
llama.cpp
chunk128
vLLM 0.24
legacy
vLLM 0.31
SGLang 0.5.19
singleton guard / BF16 SSM
27B / one RTX4090 / C4 / 0.5 short RPS / 24K short history / 128K long.
Different quantization and cache capacities; these are deployment observations, not isolated scheduler rankings.