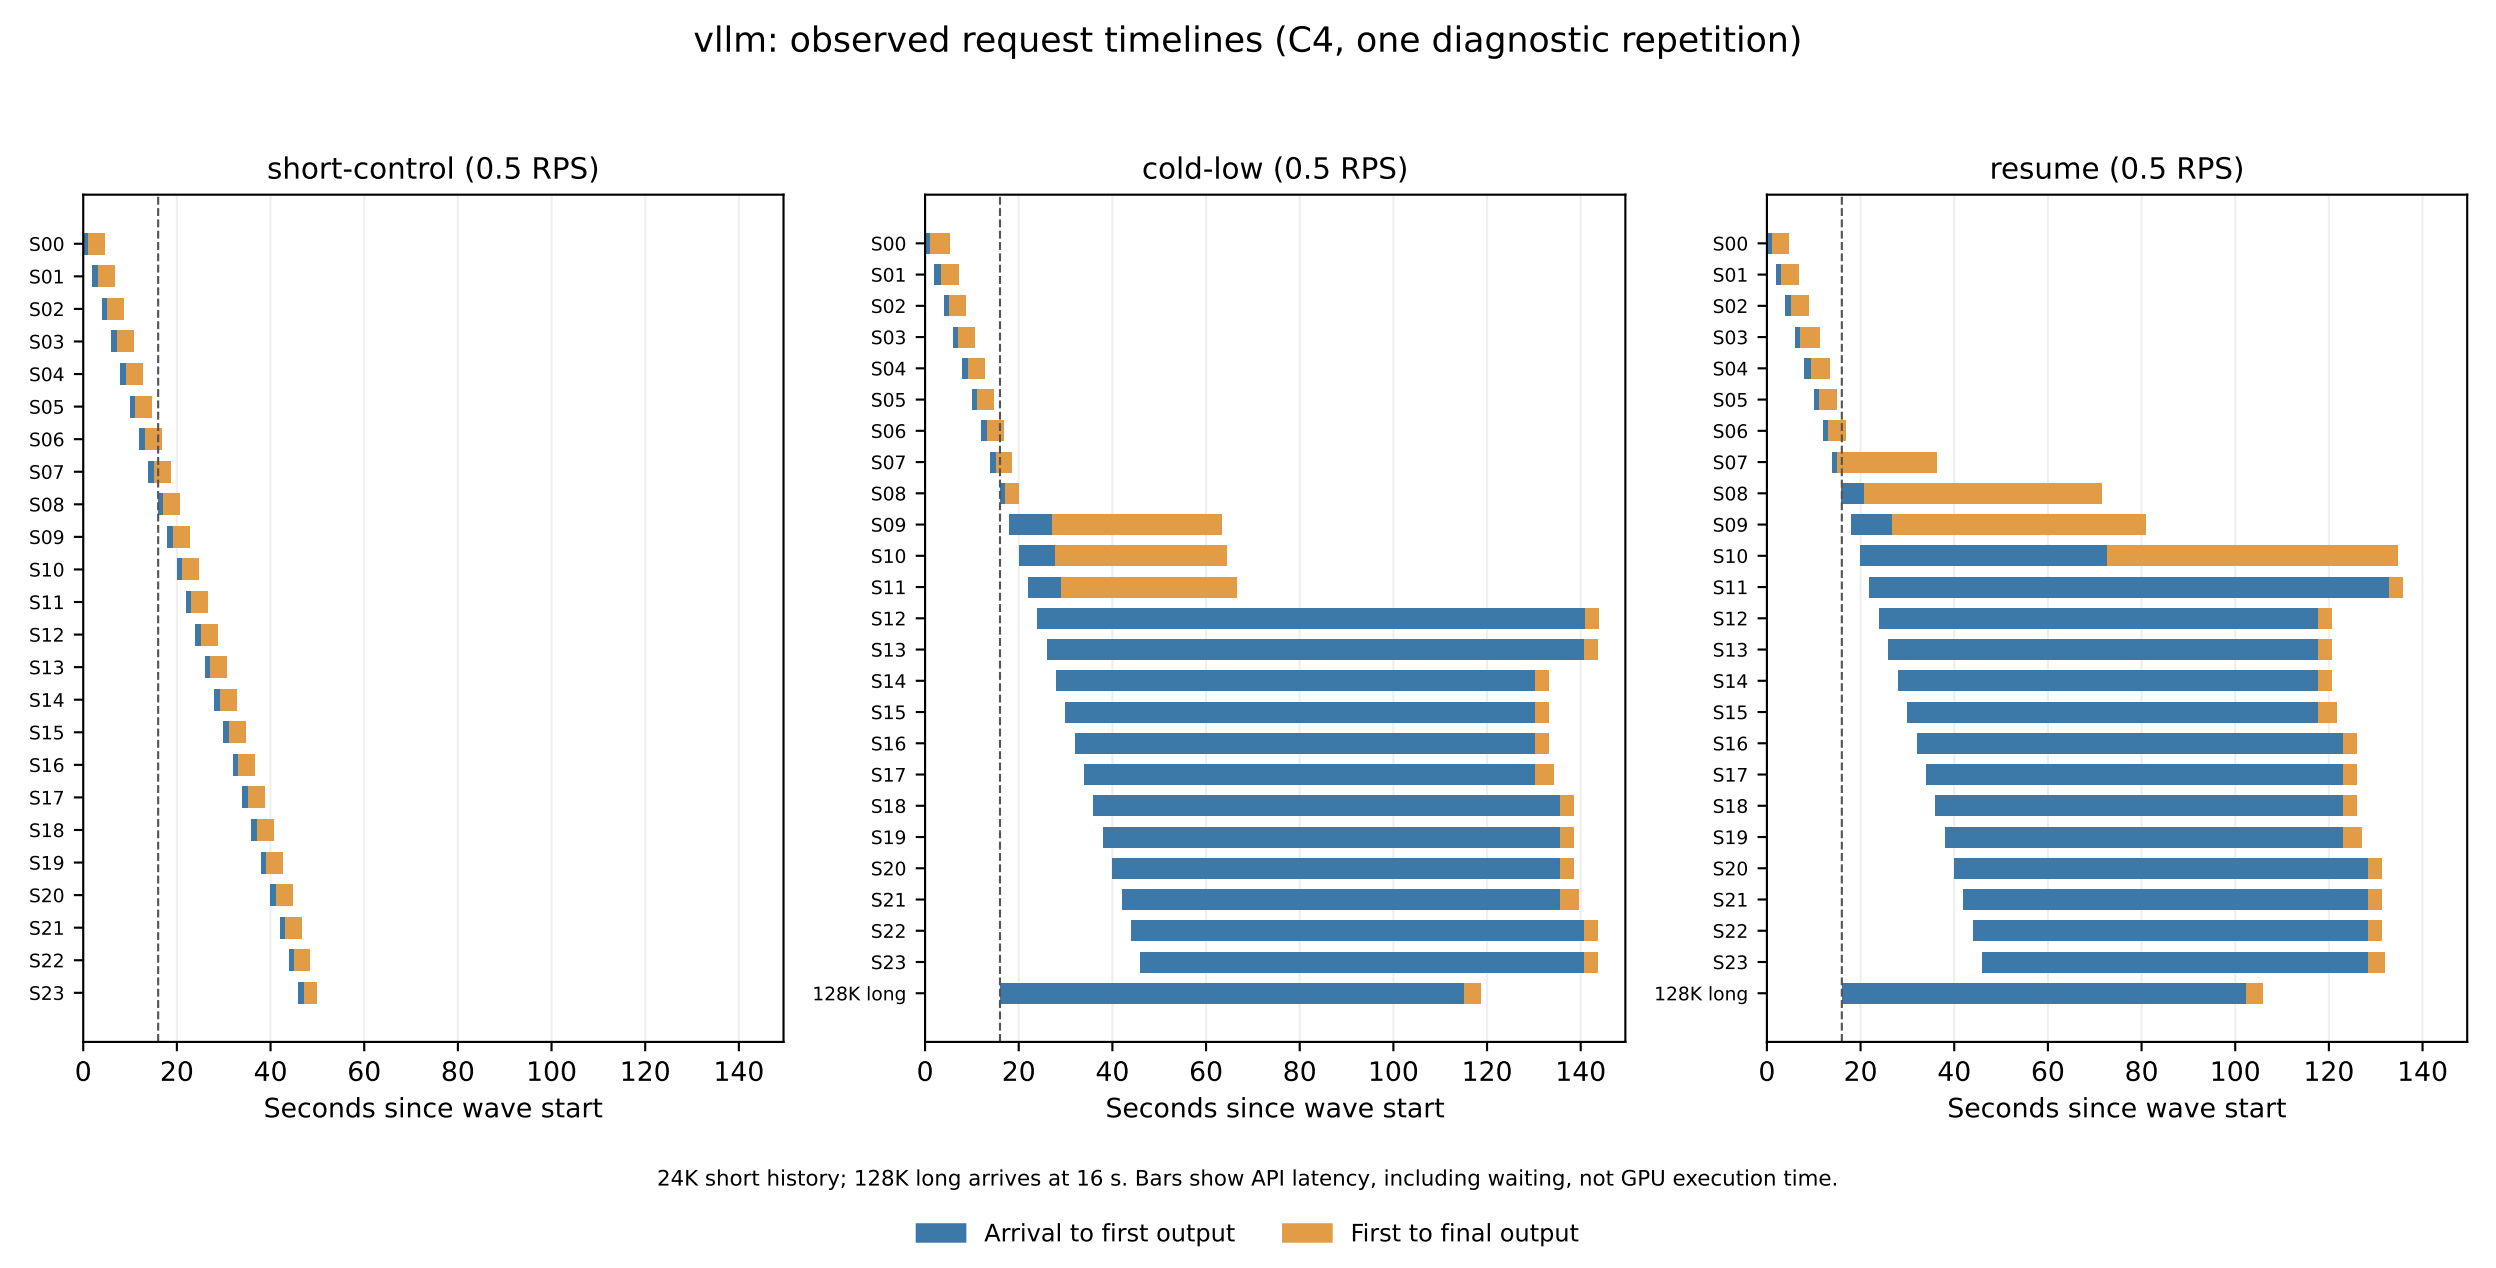
vllm: observed request timelines (C4, one diagnostic repetition)
short-control (0.5 RPS)
cold-low (0.5 RPS)
resume (0.5 RPS)
S00
S01
S02
S03
S04
S05
S06
S07
S08
S09
S10
S11
S12
S13
S14
S15
S16
S17
S18
S19
S20
S21
S22
S23
128K long
0
20
40
60
80
100
120
140
Seconds since wave start
24K short history; 128K long arrives at 16 s. Bars show API latency, including waiting, not GPU execution time.
Arrival to first output
First to final output